

**Task 1: Understanding Naive Bayes and K-nearest neighbors**

**1a: Manually calculate prediction using the Naive Bayes Model and K nearest neighbor, K=2; Euclidean Distance for the test example for the following example:**

- Use any random combination to test/report your probability

ID	Contains Link	Contains Money Words	Length	Class
1	Yes	Yes	Long	Spam
2	No	No	Short	Ham
3	Yes	No	Long	Spam
4	No	Yes	Short	Spam
5	Yes	Yes	Short	Spam
6	No	No	Long	Ham
7	Yes	No	Short	Ham
8	No	Yes	Long	Spam
9	Yes	Yes	Long	Spam
10	No	No	Short	Ham

⇒ Naive Bayes

- From table,
 
$$P(\text{Spam}) = \frac{6}{10} \Rightarrow 0.6$$

$$P(\text{Ham}) = \frac{4}{10} \Rightarrow 0.4$$
- Conditional Probabilities
  - For Spam,
 
$$P(\text{Yes}|\text{Spam}) \Rightarrow \text{Contains link} \Rightarrow \frac{4}{6} \Rightarrow 0.67$$

$$P(\text{No}|\text{Spam}) \Rightarrow \text{Contains money} \Rightarrow \frac{5}{6} \Rightarrow 0.833$$

$$P(\text{Long}|\text{Spam}) \Rightarrow \text{Long} \Rightarrow \frac{4}{6} \Rightarrow 0.67$$
  - For Ham,
 
$$P(\text{Contains link}) \Rightarrow (\text{No}/\text{Ham}) = \frac{3}{4} = 0.75$$

$$P(\text{No}/\text{Ham}) \Rightarrow \text{Contains Money} = \frac{4}{4} = 1$$

$$P(\text{Long}/\text{Ham}) \Rightarrow \text{Length} = \frac{1}{4} = 0.25$$
- Posterior Probability
  - For Spam =  $0.6 \times 0.67 \times 0.833 \times 0.67$   

$$= 0.224$$
  - For Ham =  $0.4 \times 0.75 \times 1 \times 0.25 = 0.075 //$

Thus,

$$P(\text{Spam}|\text{Features}) > P(\text{Ham}|\text{Features})$$

The class is SPAM

(1a)  $\Rightarrow$  KNN

ID	Contains Link	Contains Money words	length	Class
1	Yes $\rightarrow 1$	Yes $\rightarrow 1$	long $\rightarrow 1$	Spam $\Rightarrow \sqrt{1^2+1^2+1^2} \Rightarrow \sqrt{3} \Rightarrow 1.7^*$
2	No $\rightarrow 0$	No $\rightarrow 0$	short $\rightarrow 0$	Ham $\Rightarrow \sqrt{0} \Rightarrow 0$
3	Yes $\rightarrow 1$	No $\rightarrow 0$	long $\rightarrow 1$	Spam $\Rightarrow \sqrt{2} \Rightarrow 1.4^*$
4	No $\rightarrow 0$	Yes $\rightarrow 1$	short $\rightarrow 0$	Spam $\Rightarrow \sqrt{1} \Rightarrow 1$
5	Yes $\rightarrow 1$	Yes $\rightarrow 1$	short $\rightarrow 0$	Spam $\Rightarrow \sqrt{2} \Rightarrow 1.4^*$
6	No $\rightarrow 0$	No $\rightarrow 0$	long $\rightarrow 1$	Ham $\Rightarrow \sqrt{1} \Rightarrow 1$
7	Yes $\rightarrow 1$	No $\rightarrow 0$	short $\rightarrow 0$	Ham $\Rightarrow \sqrt{1} \Rightarrow 1$
8	No $\rightarrow 0$	Yes $\rightarrow 1$	long $\rightarrow 1$	Spam $\Rightarrow \sqrt{2} \Rightarrow 1.4^*$
9	Yes $\rightarrow 1$	Yes $\rightarrow 1$	long $\rightarrow 1$	Spam $\Rightarrow \sqrt{3} \Rightarrow 1.7^*$
10	No $\rightarrow 0$	No $\rightarrow 0$	short $\rightarrow 0$	Ham $\Rightarrow \sqrt{0} \Rightarrow 0$

$\Rightarrow$  Ordering the values:-

ID	value	Class
1	1.7	Spam
9	1.7	Spam
3	1.4	Spam
5	1.4	Spam
8	1.4	Spam
4	1	Spam
6	1	Ham
7	1	Ham
10	0	Ham
2	0	Ham

From the computation, consider the range 1.4. Any class with value above 1.4 should be spam. Any class with below 1.4 should be ham.

$\Rightarrow$  Now consider test example,

ID	Contains Link	Contains Money words	length	Class	value
11	Yes $\rightarrow 1$	No $\rightarrow 0$	short $\rightarrow 0$	?	$\sqrt{1} \Rightarrow 1 < 1.4 \Rightarrow \text{Ham}$
12	No $\rightarrow 0$	Yes $\rightarrow 1$	long $\rightarrow 1$	?	$\sqrt{2} \Rightarrow 1.4 = 1.4 \Rightarrow \text{Spam}$

1b: write code (with AI assistant) to build a naive Bayes and KNN classifier. You can use the hamspam.csv to test it out.



Q Commands + Code + Text



12s

```
[11] from google.colab import files
      uploaded = files.upload()
```



Choose Files hamspam.csv.csv



- **hamspam.csv.csv**(text/csv) - 20909 bytes, last modified: 3/3/2025 - 100% done  
Saving hamspam.csv.csv to hamspam.csv.csv



0s



# KNN

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# Load the dataset
df = pd.read_csv('/content/hamspam.csv.csv')

# Encode categorical columns if any
le = LabelEncoder()
for col in X.columns:
    if X[col].dtype == 'object':
        X[col] = le.fit_transform(X[col])

# Split dataset into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create KNN classifier with k=2
knn = KNeighborsClassifier(n_neighbors=2)

# Fit the model
knn.fit(X_train, y_train)

# Make predictions
y_pred = knn.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
print("Classification Report:\n", classification_report(y_test, y_pred)) # Fixed space issue
```



Q Commands + Code + Text

```
[36] print("Classification Report:\n", classification_report(y_test, y_pred)) # Fixed space issue
```

Accuracy: 50.50%

Classification Report:

	precision	recall	f1-score	support
Ham	0.51	0.72	0.60	103
Spam	0.48	0.28	0.35	97
accuracy			0.51	200
macro avg	0.50	0.50	0.48	200
weighted avg	0.50	0.51	0.48	200

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# Load the dataset
df = pd.read_csv('/content/hamspam.csv.csv')

# Assuming the last column is the target and others are features
X = dataset.iloc[:, :-1] # Features
y = dataset.iloc[:, -1] # Target

# Encode categorical columns if any
le = LabelEncoder()
for col in X.columns:
    if X[col].dtype == 'object':
        X[col] = le.fit_transform(X[col])

# Split dataset into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create Naive Bayes classifier
nb = GaussianNB()

# Fit the model
nb.fit(X_train, y_train)

# Make predictions
y_pred = nb.predict(X_test)
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
print("Classification Report:\n", classification_report(y_test, y_pred)) # Fixed non-breaking space
```

Accuracy: 55.00%

Classification Report:

	precision	recall	f1-score	support
Ham	0.54	0.80	0.65	103
Spam	0.57	0.29	0.38	97
accuracy			0.55	200
macro avg	0.56	0.54	0.51	200
weighted avg	0.56	0.55	0.52	200

## **Task2: Understanding ROC and AUC**

2a: Create a ROC (with AI assistant/Excel ) (**Refer to** roc\_data.csv)

Step1: Given the threshold (0.95,0.90,0.85,0.80,0.75,0.70), derive True Positive and False Positive

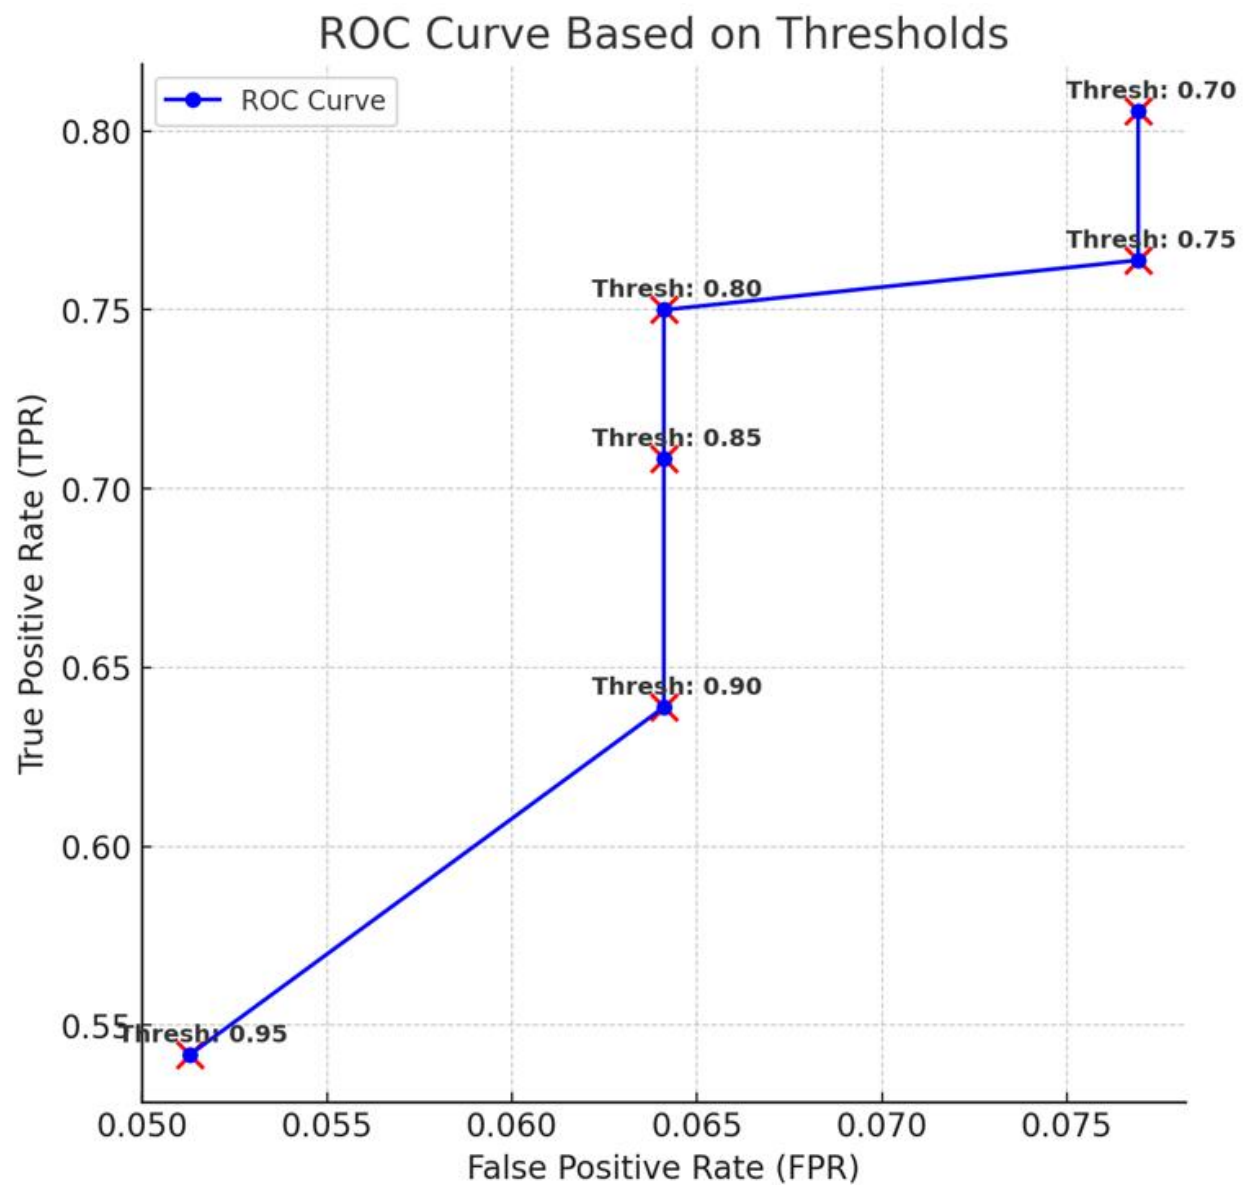
Step2: Calculate the True Positive Rate (TPR) and False Positive Rate (FPR), enter the values into the sheet

Step3: plot the set points (FRP, TPR) on the ROC diagram

a]

Threshold	TP	FP	FN	TN	TPR	FPR
0.95	39	4	33	74	0.541	0.0512
0.9	46	5	26	73	0.638	0.0641
0.85	51	5	21	73	0.708	0.0641
0.8	54	5	18	73	0.75	0.0641
0.75	55	6	17	72	0.7638	0.07692
0.7	58	6	14	72	0.805	0.07692

	A	B	C	D	E	F	G
1	Threshold	TP	FP	FN	TN	TPR	FPR
2	0.95	39	4	33	74	0.54167	0.05128
3	0.9	46	5	26	73	0.63889	0.0641
4	0.85	51	5	21	73	0.70833	0.0641
5	0.8	54	5	18	73	0.75	0.0641
6	0.75	55	6	17	72	0.76389	0.07692
7	0.7	58	6	14	72	0.80556	0.07692



2b. Write code (with AI assistant) to fit the model using your favorite classifier (NB, KNN, or Decision tree); using the hamspam.csv, ask to output an ROC curve and AUC score. (Hint: if you fit a decision tree, you might want to reduce max\_depth)



```

#24

#KNN ROC
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import LabelEncoder

# Assuming the last column is the target and others are features
X = dataset.iloc[:, :-1] # Features
y = dataset.iloc[:, -1] # Target

# Encode categorical columns if any
le = LabelEncoder()
for col in X.columns:
    if X[col].dtype == 'object':
        X[col] = le.fit_transform(X[col])

# Encode target variable if it's categorical
y = le.fit_transform(y) if y.dtype == 'object' else y

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train KNN classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Predict probabilities
if hasattr(knn, "predict_proba"):
    y_scores = knn.predict_proba(X_test)[:, 1]
else:
    y_scores = knn.predict(X_test) # Fallback if predict_proba is not available


# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_scores)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

# Print AUC Score
print(f'AUC Score: {roc_auc:.2f}')

```

↕



AUC Score: 0.49